

Data Aggregation and Group Operations

```
In [2]: import numpy as np
import pandas as pd
PREVIOUS_MAX_ROWS = pd.options.display.max_rows
pd.options.display.max_rows = 20
np.random.seed(12345)
import matplotlib.pyplot as plt
plt.rc('figure', figsize=(10, 6))
np.set_printoptions(precision=4, suppress=True)
```

GroupBy Mechanics

```
In [3]: df = pd.DataFrame({'key1' : ['a', 'a', 'b', 'b', 'a'],
                           'key2' : ['one', 'two', 'one', 'two', 'one'],
                           'data1' : np.random.randn(5),
                           'data2' : np.random.randn(5)})

df
```

```
Out[3]:
```

	key1	key2	data1	data2
0	a	one	-0.204708	1.393406
1	a	two	0.478943	0.092908
2	b	one	-0.519439	0.281746
3	b	two	-0.555730	0.769023
4	a	one	1.965781	1.246435

```
In [4]: grouped = df['data1'].groupby(df['key1'])
grouped
```

```
Out[4]: <pandas.core.groupby.generic.SeriesGroupBy object at 0x0000021DC3683CA0>
```

```
In [5]: grouped.mean()
```

```
Out[5]: key1
a      0.746672
b     -0.537585
Name: data1, dtype: float64
```

```
In [6]: means = df['data1'].groupby([df['key1'], df['key2']]).mean()
means
```

```
Out[6]: key1  key2
a      one    0.880536
         two    0.478943
b      one   -0.519439
         two   -0.555730
Name: data1, dtype: float64
```

```
In [7]: means.unstack()
```

Out[7]: **key2** **one** **two**

key1

a 0.880536 0.478943

b -0.519439 -0.555730

```
In [8]: states = np.array(['Ohio', 'California', 'California', 'Ohio', 'Ohio'])
years = np.array([2005, 2005, 2006, 2005, 2006])
df['data1'].groupby([states, years]).mean()
```

Out[8]:

California	2005	0.478943
	2006	-0.519439
Ohio	2005	-0.380219
	2006	1.965781

Name: data1, dtype: float64

```
In [9]: df.groupby('key1').mean()
```

Out[9]:

	data1	data2
--	--------------	--------------

key1

a 0.746672 0.910916

b -0.537585 0.525384

```
In [11]: df.groupby(['key1', 'key2']).mean()
```

Out[11]:

		data1	data2
--	--	--------------	--------------

key1 key2

a one 0.880536 1.319920

two 0.478943 0.092908

b one -0.519439 0.281746

two -0.555730 0.769023

```
In [12]: df.groupby(['key1', 'key2']).size()
```

Out[12]:

	key1	key2
a	one	2
	two	1
b	one	1
	two	1

dtype: int64

Iterating Over Groups

```
In [13]: for name, group in df.groupby('key1'):
          print(name)
          print(group)
```

```

a
  key1 key2    data1    data2
0    a  one -0.204708  1.393406
1    a  two  0.478943  0.092908
4    a  one  1.965781  1.246435
b
  key1 key2    data1    data2
2    b  one -0.519439  0.281746
3    b  two -0.555730  0.769023

```

```
In [14]: for (k1, k2), group in df.groupby(['key1', 'key2']):
          print((k1, k2))
          print(group)
```

```

('a', 'one')
  key1 key2    data1    data2
0    a  one -0.204708  1.393406
4    a  one  1.965781  1.246435
('a', 'two')
  key1 key2    data1    data2
1    a  two  0.478943  0.092908
('b', 'one')
  key1 key2    data1    data2
2    b  one -0.519439  0.281746
('b', 'two')
  key1 key2    data1    data2
3    b  two -0.55573  0.769023

```

```
In [15]: pieces = dict(list(df.groupby('key1')))
          pieces['b']
```

```
Out[15]:
```

	key1	key2	data1	data2
2	b	one	-0.519439	0.281746
3	b	two	-0.555730	0.769023

```
In [16]: df.dtypes
grouped = df.groupby(df.dtypes, axis=1)
```

```
In [17]: for dtype, group in grouped:
          print(dtype)
          print(group)
```

```

float64
  data1    data2
0 -0.204708  1.393406
1  0.478943  0.092908
2 -0.519439  0.281746
3 -0.555730  0.769023
4  1.965781  1.246435
object
  key1 key2
0    a  one
1    a  two
2    b  one
3    b  two
4    a  one

```

Selecting a Column or Subset of Columns

```
In [18]: df.groupby('key1')['data1']
          df.groupby('key1')[['data2']]
```

Out[18]: <pandas.core.groupby.generic.DataFrameGroupBy object at 0x0000021DC365DCD0>

```
In [19]: df['data1'].groupby(df['key1'])
df[['data2']].groupby(df['key1'])
```

Out[19]: <pandas.core.groupby.generic.DataFrameGroupBy object at 0x0000021DC05E9940>

```
In [20]: df.groupby(['key1', 'key2'])['data2'].mean()
```

Out[20]:

		data2	
	key1	key2	
	a	one	1.319920
		two	0.092908
	b	one	0.281746
		two	0.769023

```
In [21]: s_grouped = df.groupby(['key1', 'key2'])['data2']
s_grouped
s_grouped.mean()
```

Out[21]:

	key1	key2	
	a	one	1.319920
		two	0.092908
	b	one	0.281746
		two	0.769023

Name: data2, dtype: float64

Grouping with Dicts and Series

```
In [22]: people = pd.DataFrame(np.random.randn(5, 5),
                                columns=['a', 'b', 'c', 'd', 'e'],
                                index=['Joe', 'Steve', 'Wes', 'Jim', 'Travis'])
people.iloc[2:3, [1, 2]] = np.nan # Add a few NA values
people
```

Out[22]:

	a	b	c	d	e
Joe	1.007189	-1.296221	0.274992	0.228913	1.352917
Steve	0.886429	-2.001637	-0.371843	1.669025	-0.438570
Wes	-0.539741	NaN	NaN	-1.021228	-0.577087
Jim	0.124121	0.302614	0.523772	0.000940	1.343810
Travis	-0.713544	-0.831154	-2.370232	-1.860761	-0.860757

```
In [23]: mapping = {'a': 'red', 'b': 'red', 'c': 'blue',
                    'd': 'blue', 'e': 'red', 'f': 'orange'}
```

```
In [24]: by_column = people.groupby(mapping, axis=1)
by_column.sum()
```

```
Out[24]:
```

	blue	red
Joe	0.503905	1.063885
Steve	1.297183	-1.553778
Wes	-1.021228	-1.116829
Jim	0.524712	1.770545
Travis	-4.230992	-2.405455

```
In [25]: map_series = pd.Series(mapping)
map_series
```

```
Out[25]:
```

a	red
b	red
c	blue
d	blue
e	red
f	orange

dtype: object

```
In [26]: people.groupby(map_series, axis=1).count()
```

```
Out[26]:
```

	blue	red
Joe	2	3
Steve	2	3
Wes	1	2
Jim	2	3
Travis	2	3

```
In [28]: people.groupby(map_series, axis=1).describe()
```

```
Out[28]:
```

	count	mean	std	min	25%	50%	75%	max
a	5.0	0.152891	0.790440	-0.713544	-0.539741	0.124121	0.886429	1.007189
b	4.0	-0.956600	0.967609	-2.001637	-1.472575	-1.063687	-0.547712	0.302614
c	4.0	-0.485828	1.311756	-2.370232	-0.871440	-0.048425	0.337187	0.523772
d	5.0	-0.196622	1.336982	-1.860761	-1.021228	0.000940	0.228913	1.669025
e	5.0	0.164062	1.091776	-0.860757	-0.577087	-0.438570	1.343810	1.352917

Grouping with Functions

```
In [29]: people.groupby(len).sum()
```

```
Out[29]:
```

	a	b	c	d	e
3	0.591569	-0.993608	0.798764	-0.791374	2.119639
5	0.886429	-2.001637	-0.371843	1.669025	-0.438570
6	-0.713544	-0.831154	-2.370232	-1.860761	-0.860757

```
In [30]: people.groupby(len, axis = 1).sum()
```

```
Out[30]:
```

	1
Joe	1.567790
Steve	-0.256595
Wes	-2.138056
Jim	2.295257
Travis	-6.636447

```
In [31]: key_list = ['one', 'one', 'one', 'two', 'two']
         people.groupby([len, key_list]).min()
```

```
Out[31]:
```

		a	b	c	d	e
3	one	-0.539741	-1.296221	0.274992	-1.021228	-0.577087
	two	0.124121	0.302614	0.523772	0.000940	1.343810
5	one	0.886429	-2.001637	-0.371843	1.669025	-0.438570
6	two	-0.713544	-0.831154	-2.370232	-1.860761	-0.860757

Grouping by Index Levels

```
In [32]: columns = pd.MultiIndex.from_arrays([['US', 'US', 'US', 'JP', 'JP'],
                                             [1, 3, 5, 1, 3]],
                                             names=['cty', 'tenor'])
         hier_df = pd.DataFrame(np.random.randn(4, 5), columns=columns)
         hier_df
```

```
Out[32]:
```

	cty			US		JP
	tenor	1	3	5	1	3
0		0.560145	-1.265934	0.119827	-1.063512	0.332883
1		-2.359419	-0.199543	-1.541996	-0.970736	-1.307030
2		0.286350	0.377984	-0.753887	0.331286	1.349742
3		0.069877	0.246674	-0.011862	1.004812	1.327195

```
In [33]: hier_df.groupby(level='cty', axis=1).count()
```

```
Out[33]:
```

	cty	JP	US
0	2	3	
1	2	3	
2	2	3	
3	2	3	

Data Aggregation

In [34]: df

```
Out[34]:
```

	key1	key2	data1	data2
0	a	one	-0.204708	1.393406
1	a	two	0.478943	0.092908
2	b	one	-0.519439	0.281746
3	b	two	-0.555730	0.769023
4	a	one	1.965781	1.246435

```
In [35]: grouped = df.groupby('key1')
grouped['data1'].quantile(0.9)
```

```
Out[35]: key1
a      1.668413
b     -0.523068
Name: data1, dtype: float64
```

```
In [36]: def peak_to_peak(arr):
          return arr.max() - arr.min()
grouped.agg(peak_to_peak)
```

C:\Users\Usuário\AppData\Local\Temp\ipykernel_6776\663192242.py:3: FutureWarning: ['key2'] did not aggregate successfully. If any error is raised this will raise in a future version of pandas. Drop these columns/ops to avoid this warning.

```
grouped.agg(peak_to_peak)
```

```
Out[36]:
```

	data1	data2
key1		
a	2.170488	1.300498
b	0.036292	0.487276

```
In [37]: grouped.describe()
```

```
Out[37]:
```

	count	mean	std	min	25%	50%	75%	max	count
key1									
a	3.0	0.746672	1.109736	-0.204708	0.137118	0.478943	1.222362	1.965781	3.0 0.9
b	2.0	-0.537585	0.025662	-0.555730	-0.546657	-0.537585	-0.528512	-0.519439	2.0 0.5

Column-Wise and Multiple Function Application

```
In [38]: tips = pd.read_csv('examples/tips.csv')
tips
```

```
Out[38]:
```

	total_bill	tip	smoker	day	time	size
0	16.99	1.01	No	Sun	Dinner	2
1	10.34	1.66	No	Sun	Dinner	3
2	21.01	3.50	No	Sun	Dinner	3
3	23.68	3.31	No	Sun	Dinner	2
4	24.59	3.61	No	Sun	Dinner	4
...
239	29.03	5.92	No	Sat	Dinner	3
240	27.18	2.00	Yes	Sat	Dinner	2
241	22.67	2.00	Yes	Sat	Dinner	2
242	17.82	1.75	No	Sat	Dinner	2
243	18.78	3.00	No	Thur	Dinner	2

244 rows × 6 columns

```
In [39]: # Add tip percentage of total bill
tips['tip_pct'] = tips['tip'] / tips['total_bill']
tips[:6]
```

```
Out[39]:
```

	total_bill	tip	smoker	day	time	size	tip_pct
0	16.99	1.01	No	Sun	Dinner	2	0.059447
1	10.34	1.66	No	Sun	Dinner	3	0.160542
2	21.01	3.50	No	Sun	Dinner	3	0.166587
3	23.68	3.31	No	Sun	Dinner	2	0.139780
4	24.59	3.61	No	Sun	Dinner	4	0.146808
5	25.29	4.71	No	Sun	Dinner	4	0.186240

```
In [40]: grouped = tips.groupby(['day', 'smoker'])
```

```
In [41]: grouped_pct = grouped['tip_pct']
grouped_pct.agg('mean')
```

```
Out[41]:
```

day	smoker	
Fri	No	0.151650
	Yes	0.174783
Sat	No	0.158048
	Yes	0.147906
Sun	No	0.160113
	Yes	0.187250
Thur	No	0.160298
	Yes	0.163863

Name: tip_pct, dtype: float64

```
In [42]: grouped_pct.agg(['mean', 'std', 'peak_to_peak'])
```


Out[42]:

		mean	std	peak_to_peak
day	smoker			
Fri	No	0.151650	0.028123	0.067349
	Yes	0.174783	0.051293	0.159925
Sat	No	0.158048	0.039767	0.235193
	Yes	0.147906	0.061375	0.290095
Sun	No	0.160113	0.042347	0.193226
	Yes	0.187250	0.154134	0.644685
Thur	No	0.160298	0.038774	0.193350
	Yes	0.163863	0.039389	0.151240

In [43]: `grouped_pct.agg([('foo', 'mean'), ('bar', np.std)])`

Out[43]:

		foo	bar
day	smoker		
Fri	No	0.151650	0.028123
	Yes	0.174783	0.051293
Sat	No	0.158048	0.039767
	Yes	0.147906	0.061375
Sun	No	0.160113	0.042347
	Yes	0.187250	0.154134
Thur	No	0.160298	0.038774
	Yes	0.163863	0.039389

In [44]: `functions = ['count', 'mean', 'max']
result = grouped['tip_pct', 'total_bill'].agg(functions)
result`

C:\Users\Usuário\AppData\Local\Temp\ipykernel_6776\576261660.py:2: FutureWarning:
Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.
result = grouped['tip_pct', 'total_bill'].agg(functions)

Out[44]:

		tip_pct			total_bill		
		count	mean	max	count	mean	max
day	smoker						
Fri	No	4	0.151650	0.187735	4	18.420000	22.75
	Yes	15	0.174783	0.263480	15	16.813333	40.17
Sat	No	45	0.158048	0.291990	45	19.661778	48.33
	Yes	42	0.147906	0.325733	42	21.276667	50.81
Sun	No	57	0.160113	0.252672	57	20.506667	48.17
	Yes	19	0.187250	0.710345	19	24.120000	45.35
Thur	No	45	0.160298	0.266312	45	17.113111	41.19
	Yes	17	0.163863	0.241255	17	19.190588	43.11

In [45]: `result['tip_pct']`

Out[45]:

		count	mean	max
day	smoker			
Fri	No	4	0.151650	0.187735
	Yes	15	0.174783	0.263480
Sat	No	45	0.158048	0.291990
	Yes	42	0.147906	0.325733
Sun	No	57	0.160113	0.252672
	Yes	19	0.187250	0.710345
Thur	No	45	0.160298	0.266312
	Yes	17	0.163863	0.241255

In [46]: `ftuples = [('Durchschnitt', 'mean'), ('Abweichung', np.var)]
grouped['tip_pct', 'total_bill'].agg(ftuples)`

C:\Users\Usuário\AppData\Local\Temp\ipykernel_6776\3909980601.py:2: FutureWarning: Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.

```
grouped['tip_pct', 'total_bill'].agg(ftuples)
```

Out[46]:

		tip_pct		total_bill	
		Durchschnitt	Abweichung	Durchschnitt	Abweichung
day	smoker				
Fri	No	0.151650	0.000791	18.420000	25.596333
	Yes	0.174783	0.002631	16.813333	82.562438
Sat	No	0.158048	0.001581	19.661778	79.908965
	Yes	0.147906	0.003767	21.276667	101.387535
Sun	No	0.160113	0.001793	20.506667	66.099980
	Yes	0.187250	0.023757	24.120000	109.046044
Thur	No	0.160298	0.001503	17.113111	59.625081
	Yes	0.163863	0.001551	19.190588	69.808518

In [47]: `grouped.agg({'tip' : np.max, 'size' : 'sum'})`

Out[47]:

		tip	size
day	smoker		
Fri	No	3.50	9
	Yes	4.73	31
Sat	No	9.00	115
	Yes	10.00	104
Sun	No	6.00	167
	Yes	6.50	49
Thur	No	6.70	112
	Yes	5.00	40

In [48]: `grouped.agg({'tip_pct' : ['min', 'max', 'mean', 'std'], 'size' : 'sum'})`

Out[48]:

		tip_pct				size
		min	max	mean	std	sum
day	smoker					
Fri	No	0.120385	0.187735	0.151650	0.028123	9
	Yes	0.103555	0.263480	0.174783	0.051293	31
Sat	No	0.056797	0.291990	0.158048	0.039767	115
	Yes	0.035638	0.325733	0.147906	0.061375	104
Sun	No	0.059447	0.252672	0.160113	0.042347	167
	Yes	0.065660	0.710345	0.187250	0.154134	49
Thur	No	0.072961	0.266312	0.160298	0.038774	112
	Yes	0.090014	0.241255	0.163863	0.039389	40

Returning Aggregated Data Without Row Indexes

```
In [49]: tips.groupby(['day', 'smoker'], as_index=False).mean()
```

```
Out[49]:
```

	day	smoker	total_bill	tip	size	tip_pct
0	Fri	No	18.420000	2.812500	2.250000	0.151650
1	Fri	Yes	16.813333	2.714000	2.066667	0.174783
2	Sat	No	19.661778	3.102889	2.555556	0.158048
3	Sat	Yes	21.276667	2.875476	2.476190	0.147906
4	Sun	No	20.506667	3.167895	2.929825	0.160113
5	Sun	Yes	24.120000	3.516842	2.578947	0.187250
6	Thur	No	17.113111	2.673778	2.488889	0.160298
7	Thur	Yes	19.190588	3.030000	2.352941	0.163863

Apply: General split-apply-combine

```
In [50]: def top(df, n=5, column='tip_pct'):
          return df.sort_values(by=column)[-n:]
top(tips, n=6)
```

```
Out[50]:
```

	total_bill	tip	smoker	day	time	size	tip_pct
109	14.31	4.00	Yes	Sat	Dinner	2	0.279525
183	23.17	6.50	Yes	Sun	Dinner	4	0.280535
232	11.61	3.39	No	Sat	Dinner	2	0.291990
67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
172	7.25	5.15	Yes	Sun	Dinner	2	0.710345

```
In [51]: tips.groupby('smoker').apply(top)
```

Out[51]:

		total_bill	tip	smoker	day	time	size	tip_pct
smoker								
No	88	24.71	5.85	No	Thur	Lunch	2	0.236746
	185	20.69	5.00	No	Sun	Dinner	5	0.241663
	51	10.29	2.60	No	Sun	Dinner	2	0.252672
	149	7.51	2.00	No	Thur	Lunch	2	0.266312
	232	11.61	3.39	No	Sat	Dinner	2	0.291990
Yes	109	14.31	4.00	Yes	Sat	Dinner	2	0.279525
	183	23.17	6.50	Yes	Sun	Dinner	4	0.280535
	67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
	178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
	172	7.25	5.15	Yes	Sun	Dinner	2	0.710345

In [52]: `tips.groupby(['smoker', 'day']).apply(top, n=1, column='total_bill')`

Out[52]:

		total_bill	tip	smoker	day	time	size	tip_pct
smoker day								
No	Fri	94	22.75	3.25	No	Fri	Dinner	2 0.142857
	Sat	212	48.33	9.00	No	Sat	Dinner	4 0.186220
	Sun	156	48.17	5.00	No	Sun	Dinner	6 0.103799
	Thur	142	41.19	5.00	No	Thur	Lunch	5 0.121389
Yes	Fri	95	40.17	4.73	Yes	Fri	Dinner	4 0.117750
	Sat	170	50.81	10.00	Yes	Sat	Dinner	3 0.196812
	Sun	182	45.35	3.50	Yes	Sun	Dinner	3 0.077178
	Thur	197	43.11	5.00	Yes	Thur	Lunch	4 0.115982

In [53]: `result = tips.groupby('smoker')['tip_pct'].describe()
result`

Out[53]:

	count	mean	std	min	25%	50%	75%	max
smoker								
No	151.0	0.159328	0.039910	0.056797	0.136906	0.155625	0.185014	0.291990
Yes	93.0	0.163196	0.085119	0.035638	0.106771	0.153846	0.195059	0.710345

In [54]: `result.unstack('smoker')`

```
Out[54]:
```

	smoker	
count	No	151.000000
	Yes	93.000000
mean	No	0.159328
	Yes	0.163196
std	No	0.039910
	Yes	0.085119
min	No	0.056797
	Yes	0.035638
25%	No	0.136906
	Yes	0.106771
50%	No	0.155625
	Yes	0.153846
75%	No	0.185014
	Yes	0.195059
max	No	0.291990
	Yes	0.710345
dtype:	float64	

```
In [55]: f = lambda x: x.describe()
grouped.apply(f)
```

```
Out[55]:
```

			total_bill	tip	size	tip_pct
Fri	No	count	4.000000	4.000000	4.00	4.000000
		mean	18.420000	2.812500	2.25	0.151650
		std	5.059282	0.898494	0.50	0.028123
		min	12.460000	1.500000	2.00	0.120385
		25%	15.100000	2.625000	2.00	0.137239
	
Thur	Yes	min	10.340000	2.000000	2.00	0.090014
		25%	13.510000	2.000000	2.00	0.148038
		50%	16.470000	2.560000	2.00	0.153846
		75%	19.810000	4.000000	2.00	0.194837
		max	43.110000	5.000000	4.00	0.241255
	

64 rows × 4 columns

Suppressing the Group Keys

```
In [56]: tips.groupby('smoker', group_keys=False).apply(top)
```

```
Out[56]:
```

	total_bill	tip	smoker	day	time	size	tip_pct
88	24.71	5.85	No	Thur	Lunch	2	0.236746
185	20.69	5.00	No	Sun	Dinner	5	0.241663
51	10.29	2.60	No	Sun	Dinner	2	0.252672
149	7.51	2.00	No	Thur	Lunch	2	0.266312
232	11.61	3.39	No	Sat	Dinner	2	0.291990
109	14.31	4.00	Yes	Sat	Dinner	2	0.279525
183	23.17	6.50	Yes	Sun	Dinner	4	0.280535
67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
172	7.25	5.15	Yes	Sun	Dinner	2	0.710345

Quantile and Bucket Analysis

```
In [57]: frame = pd.DataFrame({'data1': np.random.randn(1000),
                              'data2': np.random.randn(1000)})
quartiles = pd.cut(frame.data1, 4)
quartiles[:10]
```

```
Out[57]: 0      (-1.23, 0.489]
1      (-2.956, -1.23]
2      (-1.23, 0.489]
3      (0.489, 2.208]
4      (-1.23, 0.489]
5      (0.489, 2.208]
6      (-1.23, 0.489]
7      (-1.23, 0.489]
8      (0.489, 2.208]
9      (0.489, 2.208]
Name: data1, dtype: category
Categories (4, interval[float64, right]): [(-2.956, -1.23] < (-1.23, 0.489] < (0.489, 2.208] < (2.208, 3.928]]
```

```
In [58]: def get_stats(group):
          return {'min': group.min(), 'max': group.max(),
                  'count': group.count(), 'mean': group.mean()}
grouped = frame.data2.groupby(quartiles)
grouped.apply(get_stats).unstack()
```

```
Out[58]:
```

	min	max	count	mean
data1				
(-2.956, -1.23]	-3.399312	1.670835	95.0	-0.039521
(-1.23, 0.489]	-2.989741	3.260383	598.0	-0.002051
(0.489, 2.208]	-3.745356	2.954439	297.0	0.081822
(2.208, 3.928]	-1.929776	1.765640	10.0	0.024750

```
In [59]: # Return quantile numbers
grouping = pd.qcut(frame.data1, 10, labels=False)
grouped = frame.data2.groupby(grouping)
grouped.apply(get_stats).unstack()
```

Out[59]:

	min	max	count	mean
data1				
0	-3.399312	1.670835	100.0	-0.049902
1	-1.950098	2.628441	100.0	0.030989
2	-2.925113	2.527939	100.0	-0.067179
3	-2.315555	3.260383	100.0	0.065713
4	-2.047939	2.074345	100.0	-0.111653
5	-2.989741	2.184810	100.0	0.052130
6	-2.223506	2.458842	100.0	-0.021489
7	-3.056990	2.954439	100.0	-0.026459
8	-3.745356	2.735527	100.0	0.103406
9	-2.064111	2.377020	100.0	0.220122

Example: Filling Missing Values with Group-Specific Values

```
In [60]: s = pd.Series(np.random.randn(6))
s[::2] = np.nan
s
```

```
Out[60]: 0      NaN
1    -0.125921
2      NaN
3    -0.884475
4      NaN
5     0.227290
dtype: float64
```

```
In [61]: s.fillna(s.mean())
```

```
Out[61]: 0    -0.261035
1    -0.125921
2    -0.261035
3    -0.884475
4    -0.261035
5     0.227290
dtype: float64
```

```
In [62]: states = ['Ohio', 'New York', 'Vermont', 'Florida',
                  'Oregon', 'Nevada', 'California', 'Idaho']
group_key = ['East'] * 4 + ['West'] * 4
data = pd.Series(np.random.randn(8), index=states)
data
```

```
Out[62]: Ohio      0.922264
New York  -2.153545
Vermont   -0.365757
Florida   -0.375842
Oregon     0.329939
Nevada     0.981994
California 1.105913
Idaho     -1.613716
dtype: float64
```



```
In [63]: data[['Vermont', 'Nevada', 'Idaho']] = np.nan
data
```

```
Out[63]: Ohio          0.922264
New York       -2.153545
Vermont        NaN
Florida        -0.375842
Oregon         0.329939
Nevada         NaN
California     1.105913
Idaho          NaN
dtype: float64
```

```
In [64]: data.groupby(group_key).mean()
```

```
Out[64]: East    -0.535707
West      0.717926
dtype: float64
```

```
In [65]: fill_mean = lambda g: g.fillna(g.mean())
data.groupby(group_key).apply(fill_mean)
```

```
Out[65]: Ohio          0.922264
New York       -2.153545
Vermont        -0.535707
Florida        -0.375842
Oregon         0.329939
Nevada         0.717926
California     1.105913
Idaho          0.717926
dtype: float64
```

```
In [66]: fill_values = {'East': 0.5, 'West': -1}
fill_func = lambda g: g.fillna(fill_values[g.name])
data.groupby(group_key).apply(fill_func)
```

```
Out[66]: Ohio          0.922264
New York       -2.153545
Vermont        0.500000
Florida        -0.375842
Oregon         0.329939
Nevada        -1.000000
California     1.105913
Idaho         -1.000000
dtype: float64
```

Example: Random Sampling and Permutation

```
In [68]: # Hearts, Spades, Clubs, Diamonds
suits = ['H', 'S', 'C', 'D']
card_val = (list(range(1, 11)) + [10] * 3) * 4
base_names = ['A'] + list(range(2, 11)) + ['J', 'K', 'Q']
cards = []
for suit in ['H', 'S', 'C', 'D']:
    cards.extend(str(num) + suit for num in base_names)

deck = pd.Series(card_val, index=cards)
```

```
In [69]: deck[:13]
```

```
Out[69]: AH      1
        2H      2
        3H      3
        4H      4
        5H      5
        6H      6
        7H      7
        8H      8
        9H      9
        10H     10
        JH      10
        KH      10
        QH      10
        dtype: int64
```

```
In [70]: def draw(deck, n=5):
        return deck.sample(n)
        draw(deck)
```

```
Out[70]: AD      1
        8C      8
        5H      5
        KC      10
        2C      2
        dtype: int64
```

```
In [71]: get_suit = lambda card: card[-1] # Last letter is suit
        deck.groupby(get_suit).apply(draw, n=2)
```

```
Out[71]: C  2C      2
        3C      3
        D  KD      10
        8D      8
        H  KH      10
        3H      3
        S  2S      2
        4S      4
        dtype: int64
```

```
In [72]: deck.groupby(get_suit, group_keys=False).apply(draw, n=2)
```

```
Out[72]: KC      10
        JC      10
        AD      1
        5D      5
        5H      5
        6H      6
        7S      7
        KS      10
        dtype: int64
```

Example: Group Weighted Average and Correlation

```
In [73]: df = pd.DataFrame({'category': ['a', 'a', 'a', 'a',
                                         'b', 'b', 'b', 'b'],
                           'data': np.random.randn(8),
                           'weights': np.random.rand(8)})
        df
```

Out[73]:

	category	data	weights
0	a	1.561587	0.957515
1	a	1.219984	0.347267
2	a	-0.482239	0.581362
3	a	0.315667	0.217091
4	b	-0.047852	0.894406
5	b	-0.454145	0.918564
6	b	-0.556774	0.277825
7	b	0.253321	0.955905

0	a	1.561587	0.957515
1	a	1.219984	0.347267
2	a	-0.482239	0.581362
3	a	0.315667	0.217091
4	b	-0.047852	0.894406
5	b	-0.454145	0.918564
6	b	-0.556774	0.277825
7	b	0.253321	0.955905

```
In [74]: grouped = df.groupby('category')
get_wavg = lambda g: np.average(g['data'], weights=g['weights'])
grouped.apply(get_wavg)
```

Out[74]:

category	
a	0.811643
b	-0.122262

dtype: float64

```
In [75]: close_px = pd.read_csv('examples/stock_px_2.csv', parse_dates=True,
                                index_col=0)
close_px.info()
close_px[-4:]
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2214 entries, 2003-01-02 to 2011-10-14
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0   AAPL    2214 non-null         float64
1   MSFT    2214 non-null         float64
2   XOM     2214 non-null         float64
3   SPX     2214 non-null         float64
dtypes: float64(4)
memory usage: 86.5 KB
```

Out[75]:

	AAPL	MSFT	XOM	SPX
2011-10-11	400.29	27.00	76.27	1195.54
2011-10-12	402.19	26.96	77.16	1207.25
2011-10-13	408.43	27.18	76.37	1203.66
2011-10-14	422.00	27.27	78.11	1224.58

	AAPL	MSFT	XOM	SPX
2011-10-11	400.29	27.00	76.27	1195.54
2011-10-12	402.19	26.96	77.16	1207.25
2011-10-13	408.43	27.18	76.37	1203.66
2011-10-14	422.00	27.27	78.11	1224.58

```
In [76]: spx_corr = lambda x: x.corrwith(x['SPX'])
```

```
In [77]: rets = close_px.pct_change().dropna()
```

```
In [78]: get_year = lambda x: x.year
by_year = rets.groupby(get_year)
by_year.apply(spx_corr)
```

```
Out[78]:
```

	AAPL	MSFT	XOM	SPX
2003	0.541124	0.745174	0.661265	1.0
2004	0.374283	0.588531	0.557742	1.0
2005	0.467540	0.562374	0.631010	1.0
2006	0.428267	0.406126	0.518514	1.0
2007	0.508118	0.658770	0.786264	1.0
2008	0.681434	0.804626	0.828303	1.0
2009	0.707103	0.654902	0.797921	1.0
2010	0.710105	0.730118	0.839057	1.0
2011	0.691931	0.800996	0.859975	1.0

```
In [79]: by_year.apply(lambda g: g['AAPL'].corr(g['MSFT']))
```

```
Out[79]:
```

2003	0.480868
2004	0.259024
2005	0.300093
2006	0.161735
2007	0.417738
2008	0.611901
2009	0.432738
2010	0.571946
2011	0.581987

dtype: float64

Example: Group-Wise Linear Regression

```
In [80]: import statsmodels.api as sm
def regress(data, yvar, xvars):
    Y = data[yvar]
    X = data[xvars]
    X['intercept'] = 1.
    result = sm.OLS(Y, X).fit()
    return result.params
```

```
In [81]: by_year.apply(regress, 'AAPL', ['SPX'])
```

```
Out[81]:
```

	SPX	intercept
2003	1.195406	0.000710
2004	1.363463	0.004201
2005	1.766415	0.003246
2006	1.645496	0.000080
2007	1.198761	0.003438
2008	0.968016	-0.001110
2009	0.879103	0.002954
2010	1.052608	0.001261
2011	0.806605	0.001514

Pivot Tables and Cross-Tabulation

```
In [82]: tips.pivot_table(index=['day', 'smoker'])
```

```
Out[82]:
```

		size	tip	tip_pct	total_bill
day	smoker				
Fri	No	2.250000	2.812500	0.151650	18.420000
	Yes	2.066667	2.714000	0.174783	16.813333
Sat	No	2.555556	3.102889	0.158048	19.661778
	Yes	2.476190	2.875476	0.147906	21.276667
Sun	No	2.929825	3.167895	0.160113	20.506667
	Yes	2.578947	3.516842	0.187250	24.120000
Thur	No	2.488889	2.673778	0.160298	17.113111
	Yes	2.352941	3.030000	0.163863	19.190588

```
In [83]: tips.pivot_table(['tip_pct', 'size'], index=['time', 'day'],
                           columns='smoker')
```

```
Out[83]:
```

			size		tip_pct	
	smoker		No	Yes	No	Yes
time	day					
Dinner	Fri		2.000000	2.222222	0.139622	0.165347
	Sat		2.555556	2.476190	0.158048	0.147906
	Sun		2.929825	2.578947	0.160113	0.187250
	Thur		2.000000	NaN	0.159744	NaN
Lunch	Fri		3.000000	1.833333	0.187735	0.188937
	Thur		2.500000	2.352941	0.160311	0.163863

```
In [84]: tips.pivot_table(['tip_pct', 'size'], index=['time', 'day'],
                           columns='smoker', margins=True)
```

Out[84]:

		size			tip_pct		
	smoker	No	Yes	All	No	Yes	All
time	day						
Dinner	Fri	2.000000	2.222222	2.166667	0.139622	0.165347	0.158916
	Sat	2.555556	2.476190	2.517241	0.158048	0.147906	0.153152
	Sun	2.929825	2.578947	2.842105	0.160113	0.187250	0.166897
	Thur	2.000000	NaN	2.000000	0.159744	NaN	0.159744
Lunch	Fri	3.000000	1.833333	2.000000	0.187735	0.188937	0.188765
	Thur	2.500000	2.352941	2.459016	0.160311	0.163863	0.161301
All		2.668874	2.408602	2.569672	0.159328	0.163196	0.160803

```
In [85]: tips.pivot_table('tip_pct', index=['time', 'smoker'], columns='day',
aggfunc=len, margins=True)
```

Out[85]:

		day	Fri	Sat	Sun	Thur	All
time	smoker						
Dinner	No	3.0	45.0	57.0	1.0	106	
	Yes	9.0	42.0	19.0	NaN	70	
Lunch	No	1.0	NaN	NaN	44.0	45	
	Yes	6.0	NaN	NaN	17.0	23	
All		19.0	87.0	76.0	62.0	244	

```
In [87]: tips.pivot_table('tip_pct', index=['time', 'smoker'], columns='day',
aggfunc='count', margins=True)
```

Out[87]:

		day	Fri	Sat	Sun	Thur	All
time	smoker						
Dinner	No	3.0	45.0	57.0	1.0	106	
	Yes	9.0	42.0	19.0	NaN	70	
Lunch	No	1.0	NaN	NaN	44.0	45	
	Yes	6.0	NaN	NaN	17.0	23	
All		19.0	87.0	76.0	62.0	244	

```
In [88]: tips.pivot_table('tip_pct', index=['time', 'size', 'smoker'],
columns='day', aggfunc='mean', fill_value=0)
```

Out[88]:

		day	Fri	Sat	Sun	Thur
	time	size	smoker			
Dinner	1	No	0.000000	0.137931	0.000000	0.000000
			Yes	0.000000	0.325733	0.000000
	2	No	0.139622	0.162705	0.168859	0.159744
			Yes	0.171297	0.148668	0.207893
	3	No	0.000000	0.154661	0.152663	0.000000
			Yes	0.000000	0.000000	0.000000
Lunch	3	Yes	0.000000	0.000000	0.000000	0.204952
	4	No	0.000000	0.000000	0.000000	0.138919
			Yes	0.000000	0.000000	0.155410
	5	No	0.000000	0.000000	0.000000	0.121389
			Yes	0.000000	0.000000	0.121389
	6	No	0.000000	0.000000	0.000000	0.173706

21 rows × 4 columns

Cross-Tabulations: Crosstab

```
In [89]: from io import StringIO
data = """\
Sample  Nationality  Handedness
1  USA  Right-handed
2  Japan  Left-handed
3  USA  Right-handed
4  Japan  Right-handed
5  Japan  Left-handed
6  Japan  Right-handed
7  USA  Right-handed
8  USA  Left-handed
9  Japan  Right-handed
10 USA  Right-handed"""
data = pd.read_table(StringIO(data), sep='\s+')
```

In [90]: data

Out[90]:

	Sample	Nationality	Handedness
0	1	USA	Right-handed
1	2	Japan	Left-handed
2	3	USA	Right-handed
3	4	Japan	Right-handed
4	5	Japan	Left-handed
5	6	Japan	Right-handed
6	7	USA	Right-handed
7	8	USA	Left-handed
8	9	Japan	Right-handed
9	10	USA	Right-handed

In [91]: `pd.crosstab(data.Nationality, data.Handedness, margins=True)`

Out[91]:

Handedness	Left-handed	Right-handed	All
Nationality			
Japan	2	3	5
USA	1	4	5
All	3	7	10

In [92]: `pd.crosstab([tips.time, tips.day], tips.smoker, margins=True)`

Out[92]:

	time	day	smoker	No	Yes	All
Dinner		Fri	3	9	12	
		Sat	45	42	87	
		Sun	57	19	76	
		Thur	1	0	1	
Lunch		Fri	1	6	7	
		Thur	44	17	61	
All			151	93	244	

In [93]: `pd.options.display.max_rows = PREVIOUS_MAX_ROWS`

Conclusion