

# Introduction to Modeling Libraries

```
In [1]: import numpy as np
import pandas as pd
np.random.seed(12345)
import matplotlib.pyplot as plt
plt.rc('figure', figsize=(10, 6))
PREVIOUS_MAX_ROWS = pd.options.display.max_rows
pd.options.display.max_rows = 20
np.set_printoptions(precision=4, suppress=True)
```

## Interfacing Between pandas and Model Code

```
In [2]: import pandas as pd
import numpy as np
data = pd.DataFrame({
    'x0': [1, 2, 3, 4, 5],
    'x1': [0.01, -0.01, 0.25, -4.1, 0.],
    'y': [-1.5, 0., 3.6, 1.3, -2.]})
data
```

```
Out[2]:
```

	x0	x1	y
0	1	0.01	-1.5
1	2	-0.01	0.0
2	3	0.25	3.6
3	4	-4.10	1.3
4	5	0.00	-2.0

```
In [3]: data.columns
```

```
Out[3]: Index(['x0', 'x1', 'y'], dtype='object')
```

```
In [4]: data.values
```

```
Out[4]: array([[ 1. ,  0.01, -1.5 ],
               [ 2. , -0.01,  0. ],
               [ 3. ,  0.25,  3.6 ],
               [ 4. , -4.1 ,  1.3 ],
               [ 5. ,  0. , -2. ]])
```

```
In [5]: df2 = pd.DataFrame(data.values, columns=['one', 'two', 'three'])
df2
```

Out[5]:

	one	two	three
0	1.0	0.01	-1.5
1	2.0	-0.01	0.0
2	3.0	0.25	3.6
3	4.0	-4.10	1.3
4	5.0	0.00	-2.0

```
In [6]: df3 = data.copy()
df3['strings'] = ['a', 'b', 'c', 'd', 'e']
df3
```

Out[6]:

	x0	x1	y	strings
0	1	0.01	-1.5	a
1	2	-0.01	0.0	b
2	3	0.25	3.6	c
3	4	-4.10	1.3	d
4	5	0.00	-2.0	e

```
In [7]: df3.values
```

```
Out[7]: array([[1, 0.01, -1.5, 'a'],
              [2, -0.01, 0.0, 'b'],
              [3, 0.25, 3.6, 'c'],
              [4, -4.1, 1.3, 'd'],
              [5, 0.0, -2.0, 'e']], dtype=object)
```

```
In [8]: model_cols = ['x0', 'x1']
data.loc[:, model_cols].values
```

```
Out[8]: array([[ 1. ,  0.01],
               [ 2. , -0.01],
               [ 3. ,  0.25],
               [ 4. , -4.1 ],
               [ 5. ,  0.  ]])
```

```
In [9]: data['category'] = pd.Categorical(['a', 'b', 'a', 'a', 'b'],
                                         categories=['a', 'b'])
data
```

Out[9]:

	x0	x1	y	category
0	1	0.01	-1.5	a
1	2	-0.01	0.0	b
2	3	0.25	3.6	a
3	4	-4.10	1.3	a
4	5	0.00	-2.0	b

```
In [10]: dummies = pd.get_dummies(data.category, prefix='category')
data_with_dummies = data.drop('category', axis=1).join(dummies)
data_with_dummies
```

```
Out[10]:
```

	x0	x1	y	category_a	category_b
0	1	0.01	-1.5	1	0
1	2	-0.01	0.0	0	1
2	3	0.25	3.6	1	0
3	4	-4.10	1.3	1	0
4	5	0.00	-2.0	0	1

## Creating Model Descriptions with Patsy

$y \sim x0 + x1$

```
In [11]: data = pd.DataFrame({
    'x0': [1, 2, 3, 4, 5],
    'x1': [0.01, -0.01, 0.25, -4.1, 0.],
    'y': [-1.5, 0., 3.6, 1.3, -2.]})
data
```

```
Out[11]:
```

	x0	x1	y
0	1	0.01	-1.5
1	2	-0.01	0.0
2	3	0.25	3.6
3	4	-4.10	1.3
4	5	0.00	-2.0

```
In [12]: import patsy
y, X = patsy.dmatrices('y ~ x0 + x1', data)
```

```
In [13]: y
```

```
Out[13]: DesignMatrix with shape (5, 1)
      y
-1.5
 0.0
 3.6
 1.3
-2.0
Terms:
'y' (column 0)
```

```
In [14]: X
```

```
Out[14]: DesignMatrix with shape (5, 3)
      Intercept  x0  x1
      1      1  0.01
      1      2 -0.01
      1      3  0.25
      1      4 -4.10
      1      5  0.00
Terms:
'Intercept' (column 0)
'x0' (column 1)
'x1' (column 2)
```

```
In [15]: np.asarray(y)
```

```
Out[15]: array([[ -1.5],
               [  0. ],
               [  3.6],
               [  1.3],
               [-2. ]])
```

```
In [16]: np.asarray(X)
```

```
Out[16]: array([[ 1. ,  1. ,  0.01],
               [ 1. ,  2. , -0.01],
               [ 1. ,  3. ,  0.25],
               [ 1. ,  4. , -4.1 ],
               [ 1. ,  5. ,  0. ]])
```

```
In [17]: patsy.dmatrices('y ~ x0 + x1 + 0', data)[1]
```

```
Out[17]: DesignMatrix with shape (5, 2)
          x0      x1
1      0.01
2     -0.01
3      0.25
4     -4.10
5      0.00
Terms:
  'x0' (column 0)
  'x1' (column 1)
```

```
In [18]: coef, resid, _, _ = np.linalg.lstsq(X, y)
```

C:\Users\Usuário\AppData\Local\Temp\ipykernel\_18184\2525922789.py:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)`` where M and N are the input matrix dimensions. To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old, explicitly pass `rcond=-1`.

```
coef, resid, _, _ = np.linalg.lstsq(X, y)
```

```
In [19]: coef
```

```
Out[19]: array([[ 0.3129],
               [-0.0791],
               [-0.2655]])
```

```
In [20]: coef = pd.Series(coef.squeeze(), index=X.design_info.column_names)
coef
```

```
Out[20]: Intercept    0.312910
x0             -0.079106
x1             -0.265464
dtype: float64
```

## Data Transformations in Patsy Formulas

```
In [21]: y, X = patsy.dmatrices('y ~ x0 + np.log(np.abs(x1) + 1)', data)
X
```

```
Out[21]: DesignMatrix with shape (5, 3)
      Intercept  x0  np.log(np.abs(x1) + 1)
      1      1      0.00995
      1      2      0.00995
      1      3      0.22314
      1      4      1.62924
      1      5      0.00000

Terms:
  'Intercept' (column 0)
  'x0' (column 1)
  'np.log(np.abs(x1) + 1)' (column 2)
```

```
In [22]: y, X = patsy.dmatrices('y ~ standardize(x0) + center(x1)', data)
X
```

```
Out[22]: DesignMatrix with shape (5, 3)
      Intercept  standardize(x0)  center(x1)
      1      -1.41421      0.78
      1      -0.70711      0.76
      1      0.00000      1.02
      1      0.70711      -3.33
      1      1.41421      0.77

Terms:
  'Intercept' (column 0)
  'standardize(x0)' (column 1)
  'center(x1)' (column 2)
```

```
In [23]: new_data = pd.DataFrame({
      'x0': [6, 7, 8, 9],
      'x1': [3.1, -0.5, 0, 2.3],
      'y': [1, 2, 3, 4]})
new_X = patsy.build_design_matrices([X.design_info], new_data)
new_X
```

```
Out[23]: [DesignMatrix with shape (4, 3)
      Intercept  standardize(x0)  center(x1)
      1      2.12132      3.87
      1      2.82843      0.27
      1      3.53553      0.77
      1      4.24264      3.07

Terms:
  'Intercept' (column 0)
  'standardize(x0)' (column 1)
  'center(x1)' (column 2)]
```

```
In [24]: y, X = patsy.dmatrices('y ~ I(x0 + x1)', data)
X
```

```
Out[24]: DesignMatrix with shape (5, 2)
      Intercept  I(x0 + x1)
      1      1.01
      1      1.99
      1      3.25
      1      -0.10
      1      5.00

Terms:
  'Intercept' (column 0)
  'I(x0 + x1)' (column 1)
```

## Categorical Data and Patsy

```
In [25]: data = pd.DataFrame({
      'key1': ['a', 'a', 'b', 'b', 'a', 'b', 'a', 'b'],
      'key2': [0, 1, 0, 1, 0, 1, 0, 0],
```

```
'v1': [1, 2, 3, 4, 5, 6, 7, 8],
      'v2': [-1, 0, 2.5, -0.5, 4.0, -1.2, 0.2, -1.7]
})
y, X = patsy.dmatrices('v2 ~ key1', data)
X
```

Out[25]: DesignMatrix with shape (8, 2)

Intercept	key1[T.b]
1	0
1	0
1	1
1	1
1	0
1	1
1	0
1	1

Terms:

'Intercept' (column 0)  
'key1' (column 1)

```
In [26]: y, X = patsy.dmatrices('v2 ~ key1 + 0', data)
X
```

Out[26]: DesignMatrix with shape (8, 2)

key1[a]	key1[b]
1	0
1	0
0	1
0	1
1	0
0	1
1	0
0	1

Terms:

'key1' (columns 0:2)

```
In [27]: y, X = patsy.dmatrices('v2 ~ C(key2)', data)
X
```

Out[27]: DesignMatrix with shape (8, 2)

Intercept	C(key2)[T.1]
1	0
1	1
1	0
1	1
1	0
1	1
1	0
1	0

Terms:

'Intercept' (column 0)  
'C(key2)' (column 1)

```
In [28]: data['key2'] = data['key2'].map({0: 'zero', 1: 'one'})
data
```

Out[28]:

	key1	key2	v1	v2
0	a	zero	1	-1.0
1	a	one	2	0.0
2	b	zero	3	2.5
3	b	one	4	-0.5
4	a	zero	5	4.0
5	b	one	6	-1.2
6	a	zero	7	0.2
7	b	zero	8	-1.7

In [29]: `y, X = patsy.dmatrices('v2 ~ key1 + key2', data)`  
X

Out[29]: DesignMatrix with shape (8, 3)

Intercept	key1[T.b]	key2[T.zero]
1	0	1
1	0	0
1	1	1
1	1	0
1	0	1
1	1	0
1	0	1
1	1	1

Terms:

- 'Intercept' (column 0)
- 'key1' (column 1)
- 'key2' (column 2)

In [30]: `y, X = patsy.dmatrices('v2 ~ key1 + key2 + key1:key2', data)`  
X

Out[30]: DesignMatrix with shape (8, 4)

Intercept	key1[T.b]	key2[T.zero]	key1[T.b]:key2[T.zero]
1	0	1	0
1	0	0	0
1	1	1	1
1	1	0	0
1	0	1	0
1	1	0	0
1	0	1	0
1	1	1	1

Terms:

- 'Intercept' (column 0)
- 'key1' (column 1)
- 'key2' (column 2)
- 'key1:key2' (column 3)

## Introduction to statsmodels

### Estimating Linear Models

In [31]: `import statsmodels.api as sm`  
`import statsmodels.formula.api as smf`

```
In [32]: def dnorm(mean, variance, size=1):  
        if isinstance(size, int):  
            size = size,  
            return mean + np.sqrt(variance) * np.random.randn(*size)  
  
        # For reproducibility  
        np.random.seed(12345)  
  
        N = 100  
        X = np.c_[dnorm(0, 0.4, size=N),  
                  dnorm(0, 0.6, size=N),  
                  dnorm(0, 0.2, size=N)]  
        eps = dnorm(0, 0.1, size=N)  
        beta = [0.1, 0.3, 0.5]  
  
        y = np.dot(X, beta) + eps
```

```
In [33]: X[:5]
```

```
Out[33]: array([[ -0.1295,  -1.2128,   0.5042],  
                [  0.3029,  -0.4357,  -0.2542],  
                [-0.3285,  -0.0253,   0.1384],  
                [-0.3515,  -0.7196,  -0.2582],  
                [ 1.2433,  -0.3738,  -0.5226]])
```

```
In [34]: y[:5]
```

```
Out[34]: array([ 0.4279, -0.6735, -0.0909, -0.4895, -0.1289])
```

```
In [ ]: X_model = sm.add_constant(X)  
        X_model[:5]
```

```
In [35]: model = sm.OLS(y, X)
```

```
In [36]: results = model.fit()  
        results.params
```

```
Out[36]: array([0.1783, 0.223 , 0.501 ])
```

```
In [37]: print(results.summary())
```



## OLS Regression Results

```

=====
=====
Dep. Variable:          y    R-squared (uncentered):
0.430
Model:                OLS    Adj. R-squared (uncentered):
0.413
Method:              Least Squares    F-statistic:
24.42
Date:                Tue, 10 Sep 2024    Prob (F-statistic):          7.4
4e-12
Time:                22:40:14    Log-Likelihood:          -3
4.305
No. Observations:      100    AIC:
74.61
Df Residuals:          97    BIC:
82.42
Df Model:              3
Covariance Type:      nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
x1	0.1783	0.053	3.364	0.001	0.073	0.283
x2	0.2230	0.046	4.818	0.000	0.131	0.315
x3	0.5010	0.080	6.237	0.000	0.342	0.660

```

=====
=====
Omnibus:              4.662    Durbin-Watson:          2.201
Prob(Omnibus):        0.097    Jarque-Bera (JB):      4.098
Skew:                 0.481    Prob(JB):              0.129
Kurtosis:             3.243    Cond. No.              1.74
=====
=====

```

## Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

In [38]: data = pd.DataFrame(X, columns=['col0', 'col1', 'col2'])
data['y'] = y
data[:5]

```

```

Out[38]:
   col0  col1  col2  y
0 -0.129468 -1.212753  0.504225  0.427863
1  0.302910 -0.435742 -0.254180 -0.673480
2 -0.328522 -0.025302  0.138351 -0.090878
3 -0.351475 -0.719605 -0.258215 -0.489494
4  1.243269 -0.373799 -0.522629 -0.128941

```

```

In [39]: results = smf.ols('y ~ col0 + col1 + col2', data=data).fit()
results.params

```

```

Out[39]:
Intercept    0.033559
col0         0.176149
col1         0.224826
col2         0.514808
dtype: float64

```

```

In [40]: results.tvalues

```

```
Out[40]: Intercept    0.952188  
col0      3.319754  
col1      4.850730  
col2      6.303971  
dtype: float64
```

```
In [41]: results.predict(data[:5])
```

```
Out[41]: 0    -0.002327  
1    -0.141904  
2     0.041226  
3    -0.323070  
4    -0.100535  
dtype: float64
```

## Estimating Time Series Processes

```
In [42]: init_x = 4  
  
import random  
values = [init_x, init_x]  
N = 1000  
  
b0 = 0.8  
b1 = -0.4  
noise = dnorm(0, 0.1, N)  
for i in range(N):  
    new_x = values[-1] * b0 + values[-2] * b1 + noise[i]  
    values.append(new_x)
```

```
In [49]: MAXLAGS = 5  
model = sm.tsa.AutoReg(values, lags=MAXLAGS)  
results = model.fit()
```

```
In [50]: results.params
```

```
Out[50]: array([-0.0062,  0.7845, -0.4085, -0.0136,  0.015 ,  0.0143])
```

## Introduction to scikit-learn

```
In [51]: train = pd.read_csv('datasets/titanic/train.csv')  
test = pd.read_csv('datasets/titanic/test.csv')  
train[:4]
```

Out[51]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123

In [52]:

test[:4]

Out[52]:

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	892	3	Kelly, Mr. James	male	34.5	0	0	330911	7.8292	NaN	Q
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	363272	7.0000	NaN	S
2	894	2	Myles, Mr. Thomas Francis	male	62.0	0	0	240276	9.6875	NaN	Q
3	895	3	Wirz, Mr. Albert	male	27.0	0	0	315154	8.6625	NaN	S

In [53]:

train.isnull().sum()

Out[53]:

PassengerId0  
Survived0  
Pclass0  
Name0  
Sex0  
Age177  
SibSp0  
Parch0  
Ticket0  
Fare0  
Cabin687  
Embarked2  
dtype: int64

```
In [54]: test.isnull().sum()
```

```
Out[54]: PassengerId      0
Pclass      0
Name        0
Sex         0
Age        86
SibSp       0
Parch       0
Ticket      0
Fare        1
Cabin      327
Embarked    0
dtype: int64
```

```
In [55]: impute_value = train['Age'].median()
train['Age'] = train['Age'].fillna(impute_value)
test['Age'] = test['Age'].fillna(impute_value)
```

```
In [56]: train.isnull().sum()
```

```
Out[56]: PassengerId      0
Survived      0
Pclass      0
Name        0
Sex         0
Age         0
SibSp       0
Parch       0
Ticket      0
Fare        0
Cabin      687
Embarked     2
dtype: int64
```

```
In [57]: test.isnull().sum()
```

```
Out[57]: PassengerId      0
Pclass      0
Name        0
Sex         0
Age         0
SibSp       0
Parch       0
Ticket      0
Fare        1
Cabin      327
Embarked    0
dtype: int64
```

```
In [58]: train['IsFemale'] = (train['Sex'] == 'female').astype(int)
test['IsFemale'] = (test['Sex'] == 'female').astype(int)
```

```
In [59]: predictors = ['Pclass', 'IsFemale', 'Age']
X_train = train[predictors].values
X_test = test[predictors].values
y_train = train['Survived'].values
X_train[:5]
```

```
Out[59]: array([[ 3.,  0., 22.],
               [ 1.,  1., 38.],
               [ 3.,  1., 26.],
               [ 1.,  1., 35.],
               [ 3.,  0., 35.]])
```

```
In [60]: y_train[:5]
```

```
Out[60]: array([0, 1, 1, 1, 0], dtype=int64)
```

```
In [61]: from sklearn.linear_model import LogisticRegression
         model = LogisticRegression()
```

```
In [62]: model.fit(X_train, y_train)
```

```
Out[62]: LogisticRegression()
```

```
In [63]: y_predict = model.predict(X_test)
         y_predict[:10]
```

```
Out[63]: array([0, 0, 0, 0, 1, 0, 1, 0, 1, 0], dtype=int64)
```

```
(y_true == y_predict).mean()
```

```
In [65]: from sklearn.linear_model import LogisticRegressionCV
         model_cv = LogisticRegressionCV(cv=10)
         model_cv.fit(X_train, y_train)
```

```
Out[65]: LogisticRegressionCV(cv=10)
```

```
In [66]: from sklearn.model_selection import cross_val_score
         model = LogisticRegression(C=10)
         scores = cross_val_score(model, X_train, y_train, cv=4)
         scores
```

```
Out[66]: array([0.7758, 0.7982, 0.7758, 0.7883])
```

## Continuing Your Education

```
In [ ]: pd.options.display.max_rows = PREVIOUS_MAX_ROWS
```